

SWE3004 Operating Systems, Spring 2025

Project 1. System Call

TA)

Gwanjong Park

Yunseong Shin



Project Plan

- Total 6 projects

- ~~0) Booting xv6 operating system~~

- 1) System call

- 2) CPU scheduling

- 3) Virtual memory

- 4) Page replacement

- 5) File systems



Project 1: Make system calls

- Goal: make five new system calls (getnice, setnice, ps, meminfo, waitpid)
- Synopsis [개요](#)
 - `int getnice(int pid);`
 - `int setnice(int pid, int value);`
 - `void ps(int pid);`
 - `uint64 meminfo(void);`
 - `int waitpid(int pid);`



Project 1: Make system calls

- Description
 - The *getnice* function obtains the nice value of a process.
 - The *setnice* function sets the nice value of a process.
 - The **default nice value is 20**. Lower nice values cause more favorable scheduling.
 - It will be necessary to implement the nice value before creating the system call.
 - **The range of valid nice value is 0~39**



Project 1: Make system calls

- Description (cont'd)

- In kernel, the **ps** system call prints out process(s)'s information, which includes name, pid, state and priority(nice value) of each process.
- If the pid is 0, print out all processes' information.
- Otherwise, print out corresponding process's information.
- If there is no process corresponding to the pid, print out nothing.

```
name    pid    state    priority
init     1    SLEEPING    20
sh       2    SLEEPING    20
ps       3    RUNNING    20
```

- **meminfo** prints available memory in bytes.
- **waitpid** suspends execution until the specified process terminates.



Project 1: Make system calls

- Return values
 - **getnice** : Return the nice value of target process on success. Return -1 if there is no process corresponding to the pid.
 - **setnice** : Return 0 on success. Return -1 if there is no process corresponding to the pid or the nice value is invalid.
 - **ps** : No return value. Prints the process list directly to the console.
 - **meminfo**: Returns the amount of free memory (in bytes) available in the system.
 - **waitpid**: Returns 0 when the specified process terminates successfully. Returns -1 if the process does not exist or if the calling process does not have permission to wait for it.



How to define system call

1. Add your syscall to user/usys.pl

```
entry("fork");  
entry("exit");  
entry("wait");  
entry("pipe");
```

2. Add syscall number to kernel/syscall.h

```
// System call numbers  
#define SYS_fork 1  
#define SYS_exit 2  
#define SYS_wait 3  
#define SYS_pipe 4
```

3. Add extern and syscall element in kernel/syscall.c

```
// Prototypes for the functions that handle system calls.  
extern uint64 sys_fork(void);  
extern uint64 sys_exit(void);  
extern uint64 sys_wait(void);  
extern uint64 sys_pipe(void);
```

```
// An array mapping syscall numbers from syscall.h  
// to the function that handles the system call.  
static uint64 (*syscalls[])(void) = {  
[SYS_fork] sys_fork,  
[SYS_exit] sys_exit,  
[SYS_wait] sys_wait,
```



How to define system call

4. Add a sys_ function to kernel/sysproc.c

```
uint64
sys_fork(void)
{
    return fork();
}
```

5. Add a function that performs an action to kernel/proc.c

```
// Create a new process, copying the parent.
// Sets up child kernel stack to return as if from fork() system call.
int
fork(void)
{
    int i, pid;
    struct proc *np;
    struct proc *p = myproc();

    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }

    // Copy user memory from parent to child.
```



How to define system call

6. Add a definition to kernel/defs.h and user/user.h

```
// proc.c
int      cpuid(void);
void     exit(int);
int      fork(void);
```

```
// system calls
int fork(void);
```



How to test your system call

```
#include "types.h"
#include "user.h"
#include "stat.h"

int main()
{
    int i;
    for (i=1; i<11; i++) {
        printf(1, "%d: ", i);
        if (getpname(i))
            printf(1, "Wrong pid\n");
    }
    exit();
}
```

mytest.c (just example)

“user/mytest.c” is an example code.
Create and use your own test code.

```
UPROGS=\
    $U/_cat\
    $U/_echo\
    $U/_forktest\
    $U/_grep\
    $U/_init\
    $U/_kill\
    $U/_ln\
    $U/_ls\
    $U/_mkdir\
    $U/_rm\
    $U/_sh\
    $U/_stressfs\
    $U/_usertests\
    $U/_grind\
    $U/_wc\
    $U/_zombie\
    $U/_mytest\
```

Makefile



Test with user program

```
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp 3 -nographic -global virtio-mmio.force-legacy=false -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
Student ID: 2024123456
Name: Gildong Hong
=====Your message=====
$ mytest
>>>Testing getpname:
init
>>>Testing getnice and setnice:
initial nice value: 20
nice value after setting: 10
>>>Testing ps:
name      pid    state   priority
init       1     SLEEPING   20
sh         2     SLEEPING   20
mytest     3     RUNNING    10
>>>Testing meminfo:
available memory: 133263360 bytes
>>> Testing waitpid:
wait
start1
start2
end1
done1 4 10
end2
done2 5 10
$
```



Submission

- This project is to implement only the system calls (getnice, setnice, ps, meminfo, waitpid)
 - The user program for testing is irrelevant.
- Use the ***submit*** & ***check-submission*** binary file in Ye Server
 - **\$ make clean**
 - **\$ ~swe3004/bin/submit pa1 xv6-riscv**
 - You can submit several times, and the submission history can be checked through check-submission
 - Only the last submission will be graded



Submission

- PLEASE DO NOT COPY
 - We will run inspection program on all the submissions
 - Any unannounced penalty can be given to **both students**
 - 0 points / negative points / F grade ...
- Due date: 4/1(Tue.), 23:59:59 PM
 - -25% per day for delayed submission



Questions

- If you have questions, please ask on i-campus
 - Please use the discussion board
 - Discussion board preferred over messages
- You can also visit Corporate Collaboration Center #85533
 - Please iCampus message TA before visiting
- Reading xv6 commentary will help you a lot
<https://pdos.csail.mit.edu/6.828/2023/xv6/book-riscv-rev3.pdf>



Appendix. Trap Handling Process on xv6

- How RISC-V Handles Traps and System Calls
 - System calls are executed using **ecall**.
 - The CPU jumps to the kernel's trap handler, located in **stvec**.



Appendix. Trap Handling Process on xv6

- Example : **kill** system call

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

int
main(int argc, char **argv)
{
    int i;

    if(argc < 2){
        fprintf(2, "usage: kill pid...\n");
        exit(1);
    }
    for(i=1; i<argc; i++){
        kill(atoi(argv[i]));
    }
    exit(0);
}
```

kill.c (user level)

```
// system calls
int fork(void);
int exit(int) __attribute__((noreturn));
int wait(int*);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
```

user.h



Functions defined as assembly

```
print "#include \"kernel/syscall.h\"\\n";  
sub entry {  
  my $name = shift;  
  print ".global $name\\n";  
  print "${name}:\\n";  
  print " li a7, SYS_${name}\\n";  
  print " ecall\\n";  
  print " ret\\n";  
}
```

```
entry("kill");
```

usys.pl

```
$U/usys.S : $U/usys.pl  
    perl $U/usys.pl > $U/usys.S  
  
$U/usys.o : $U/usys.S  
    $(CC) $(CFLAGS) -c -o $U/usys.o $U/usys.S
```

Makefile

```
5 #define SYS_pipe      4  
6 #define SYS_read      5  
7 #define SYS_kill      6  
8 #define SYS_exec      7  
9 #define SYS_fstat     8
```

syscall.h

subroutine entry → li a7, SYS_kill → ecall → ??? → ret



ecall → uservec → usertrap

Trap cause → scause register

```
.section trampsec
.globl trampoline
.globl usertrap
trampoline:
.align 4
.globl uservec
uservec:

# trap.c sets stvec to point here, so
# traps from user space start here,
# in supervisor mode, but with a
# user page table.
#

# save user a0 in sscratch so
# a0 can be used to get at TRAPFRAME.
cswr sscratch, a0

# each process has a separate p->trapframe memory area,
# but it's mapped to the same virtual address
# (TRAPFRAME) in every process's user page table.
li a0, TRAPFRAME

# save the user a0 in p->trapframe->a0
csrr t0, sscratch
sd t0, 112(a0)

# initialize kernel stack pointer, from p->trapframe->kernel_sp
ld sp, 8(a0)

# make tp hold the current hartid, from p->trapframe->kernel_hartid
ld tp, 32(a0)

# load the address of usertrap(), from p->trapframe->kernel_trap
ld t0, 16(a0)

# fetch the kernel page table address, from p->trapframe->kernel_satp.
ld t1, 0(a0)

# wait for any previous memory operations to complete, so that
# they use the user page table.
sfence.vma zero, zero

# install the kernel page table.
cswr satp, t1

# flush now-stale user entries from the TLB.
sfence.vma zero, zero

# jump to usertrap(), which does not return
jr t0
```

trampoline.S

SWE3004: Operating Systems, Spring 2025

```
//
// handle an interrupt, exception, or system call from user space.
// called from trampoline.S
//
void
usertrap(void)
{
    int which_dev = 0;

    if((r_sstatus() & SSTATUS_SPP) != 0)
        panic("usertrap: not from user mode");

    // send interrupts and exceptions to kerneltrap(),
    // since we're now in the kernel.
    w_stvec((uint64)kernelvec);
```

```
if(r_scause() == 8){
    // system call

    if(killed(p))
        exit(-1);

    // sepc points to the ecall instruction,
    // but we want to return to the next instruction.
    p->trapframe->epc += 4;

    // an interrupt will change sepc, scause, and sstatus,
    // so enable only now that we're done with those registers.
    intr_on();

    syscall();
```

trap.c



syscall

```
void
syscall(void)
{
    int num;
    struct proc *p = myproc();
    num = p->trapframe->a7;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        // Use num to lookup the system call function for sysnum, call it,
        // and store its return value in p->trapframe->a0
        p->trapframe->a0 = syscalls[num]();
    } else {
        printf("%d %s: unknown sys call %d\n",
            p->pid, p->name, num);
        p->trapframe->a0 = -1;
    }
}
```

```
// An array mapping syscall numbers from syscall.h
// to the function that handles the system call.
static uint64 (*syscalls[])(void) = {
    [SYS_fork]    sys_fork,
    [SYS_exit]    sys_exit,
    [SYS_wait]    sys_wait,
    [SYS_pipe]    sys_pipe,
    [SYS_read]    sys_read,
    [SYS_kill]    sys_kill,
```

syscall.c



sys_kill → kill

```
uint64
sys_kill(void)
{
    int pid;
    argint(0, &pid);
    return kill(pid);
}
```

sysproc.c

```
// Kill the process with the given pid.
// The victim won't exit until it tries to return
// to user space (see usertrap() in trap.c).
int
kill(int pid)
{
    struct proc *p;

    for(p = proc; p < &proc[NPROC]; p++){
        acquire(&p->lock);
        if(p->pid == pid){
            p->killed = 1;
            if(p->state == SLEEPING){
                // Wake process from sleep().
                p->state = RUNNABLE;
            }
            release(&p->lock);
            return 0;
        }
        release(&p->lock);
    }
    return -1;
}
```

proc.c



usertrap → usertrapret

```
syscall();
} else if((which_dev = devintr()) != 0){
    // ok
} else {
    printf("usertrap(): unexpected scause 0x%lx pid=%d\n", r_scause(), p->pid);
    printf("          sepc=0x%lx stval=0x%lx\n", r_sepc(), r_stval());
    setkilled(p);
}

if(killed(p))
    exit(-1);

// give up the CPU if this is a timer interrupt.
if(which_dev == 2)
    yield();

usertrapret();
```

trap.c, usertrap()

```
//
// return to user space
//
void
usertrapret(void)
{
    struct proc *p = myproc();

    // we're about to switch the destination of traps from
    // kerneltrap() to usertrap(), so turn off interrupts until
    // we're back in user space, where usertrap() is correct.
    intr_off();

    // send syscalls, interrupts, and exceptions to uservec in trampoline.S
    uint64 trampoline_uservec = TRAMPOLINE + (uservec - trampoline);
    w_stvec(trampoline_uservec);
```

trap.c, usertrapret()

